

Section 3 Gate / HITL Control Record

T4L Programme & Podcast Advisory Grading Model

Document status: DRAFT (procedure + control definition - screenshots / UI packs to be attached later)

Document ID: GATE-T4L-001 (working draft - align with AIMS register when assigned)

Drive folder: 04 Gate Records

System name: T4L Programme & Podcast Advisory Grading Model

Short name: T4L grading

Organisation: Data Sentinels / Transformation Leader (T4L)

System owner: Syntiche Musawu

Version: 0.1 Draft

Date created: 17 September 2026

Last updated: 17 September 2026

Linked documents:

- Impact Assessment:
docs/aims/FRM001-AI-System-Impact-Assessment-T4L-Grading-DRAFT.md
- Risk Assessment:
docs/aims/PR003-AI-Risk-Assessment-Worksheet-T4L-Grading-DRAFT.md
- Test scripts (how to prove the gate):
docs/aims/HITL-S3-Gate-Testing-Scripts-T4L-Grading-DRAFT.md
- Auto evidence pack (baseline): docs/aims/evidence/hitl/20260917-065134/
- Runnable checker: scripts/aims/hitl_gate_test.py

Filing: copy into the official AIMS gate / HITL template if provided; label DRAFT; store under Drive 04 Gate Records to T4L-grading/. Do not backdate. Attach screenshots in a later evidence annex when UI runs are complete.

0. Document control

Field	Value
Document type	AI human-in-the-loop (HITL) gate control record
Purpose	Define the gate, who operates it, what must be recorded, and how we prove it for AIMS
Scope	T4L advisory grading only
Out of scope	Oil Lab, anomaly detection, Celo, AI governance agent, LIFT scoring engine
Standard intent	Plan-Do-Check-Improve: this document is the Do control for human oversight (with Check via export / monitoring)
Honesty rule	Real dates only. Gaps listed. No invented history of mature HITL operation before August/September 2026 baseline
Screenshot status	Not required for this draft. Evidence annex reserved for later UI / CSV screenshots

1. Why this gate exists

AIMS expects human oversight to be real, not a 0.2-second click. For T4L grading, the main failure mode is automation bias: partners treating Gemini's advisory score as final without review (see awareness briefing example of ungoverned classroom AI grading).

This gate exists to ensure:

1. AI output is advisory only
2. A named human makes an accept / edit / reject decision
3. The decision is dated and attributable
4. Anti-rubber-stamp controls (criteria acknowledgement + minimum dwell) leave an audit trail
5. AI never awards programme points alone

Primary risks controlled (from PR003):

Risk ID	Title	Gate relevance
R-T4L-01	Rubber-stamping HITL	Criteria + 25s dwell + decision evidence
R-T4L-02	Unfair / biased scores	Edit / reject override path
R-T4L-03	AI treated as final authority	Advisory-only + mandatory decision before human-final outcomes
R-T4L-04	Delayed human review	Learner explain + dispute path
R-T4L-06	Model / rubric drift	HITL CSV feeds monitoring
R-T4L-07	Supplier outage	Manual human scoring still allowed

2. Gate definition

2.1 Gate name

T4L Advisory Grade Human Gate

2.2 What must pass through the gate

Any learner programme / podcast submission that receives an AI advisory grade (ai_grade.status = completed) before that estimate is treated as the partner's accepted outcome.

2.3 Gate rule (normative)

Where an AI advisory grade exists, the partner must record one of: accept, edit, or reject, after acknowledging review criteria. Accept is blocked until the partner has spent at least 25 seconds in the review drawer (anti-rubber-stamp), unless a prior decision already exists. The AI grade must not by itself change submission status to a final learner outcome or award programme points.

2.4 What the gate is not

- Not optional "nice to have" review when AI has scored
- Not a rubber-stamp button
- Not a substitute for organisational accountability
- Not approval of the Gemini supplier model itself (that sits in REG 009 / supplier controls)

3. Operating context

Item	Detail
Application	https://app.t4leader.com
Partner surface	Programme submissions (/partner/programme-submissions)
Model	Google Gemini (gemini-3-flash-preview) via grade-submission Edge Function
Data store	Supabase programme_component_submissions
Decision authority	Partner / organisation reviewer (human)
System owner	Syntiche Musawu

4. Roles and responsibilities

Role	Responsibility
Learner	Submits work; may view AI explanation; may dispute / request human review if partner has not decided
Partner reviewer	Operates the gate: reads criteria, reviews answers, accept / edit / reject, sets status / score, saves
System owner (Syntiche)	Maintains gate design, evidence fields, test script, and this control record
AIMS coordinator	Ensures gate records are filed, dated, and available for Stage 1 readiness
Management	Signs acceptability that residual risk is OK with this gate operating as described

5. HITL procedure (standard operating steps)

5.1 Happy path - partner gate

1. Learner submits eligible artefact.
2. System triggers advisory grading; ai_grade written when complete.
3. Partner opens Programme submissions.
4. Partner opens the submission review drawer.
5. Partner reviews criteria ("what good looks like") and acknowledges them.
6. Partner reviews learner answers against criteria.
7. For Accept: wait ≥ 25 seconds dwell in the drawer, then select Accept.
8. Or select Edit (enter human score different from AI if overriding) or Reject.
9. Partner sets status / notes as required and saves.
10. System persists decision fields (see Section 6).
11. Only after human save do normal partner / checklist / points rules apply - not AI alone.

5.2 Fallback - AI unavailable or error

1. If ai_grade is missing or error, partner reviews manually.
2. Partner sets score / status without Accept of an AI estimate.
3. Record remains a valid human gate action (supports R-T4L-07).

5.3 Fallback - delayed partner (learner path)

1. Learner may see AI estimate with explanation ("explain itself").
2. If no partner decision yet, learner may dispute / request human review.
3. Status moves toward human attention (e.g. needs_revision).
4. Partner then completes Section 5.1.

5.4 Prohibited behaviours

- Accepting AI without reading criteria
- Accepting within seconds to clear backlog (defeats dwell control)
- Treating AI score as final without recorded decision
- Pasting learner submission content into unapproved consumer AI tools (Tier 3 data)
- Backdating decision timestamps or evidence files

6. Evidence the gate must produce (audit fields)

Every gated decision should leave dated records. Minimum fields:

Evidence field	Meaning
ai_grade (status, score, model, graded_at, feedback)	Advisory AI output
ai_decision (accept edit reject)	Human choice
ai_decision_at	When the human decided
criteria_acknowledged / criteria_acknowledged_at	Criteria-first control
review_opened_at	When review drawer opened
review_duration_ms	Dwell / time-in-review
score / final_score / partner_score_50	Human scoring outcomes
reviewed_by / reviewer_name / reviewed_at	Attribution
status	Workflow state after human action
learner_disputed_at / learner_dispute_note	Delayed-human path (if used)

Export: Partner UI Export AI vs human CSV and/or scripts/aims/hitl_gate_test.py --auto produce machine-readable evidence for monitoring and Stage 1 packs.

Product references (implementation):

- Partner UI: src/pages/partner/ProgrammeSubmissionsPage.tsx (ACCEPT_MIN_DWELL_SEC = 25)
- HITL metrics / CSV: src/utils/programmeAiHitl.ts
- Learner explain / dispute: src/components/courses/LearnerAiExplainPanel.tsx
- Grading function: supabase/functions/grade-submission/

7. Control design summary (what "good HITL" looks like)

Control	Design	Pass condition
Advisory-only	AI does not change status / award points alone	AI-complete rows awaiting human remain non-final until partner save
Mandatory decision	accept / edit / reject when AI grade exists	ai_decision populated on gated outcomes
Criteria-first	Acknowledge criteria before deciding	criteria_acknowledged = true on decided rows
Anti-rubber-stamp	Accept blocked for 25s	review_duration_ms typically >= 20,000 on Accept path
Human override	Edit / Reject allowed	Edited/rejected samples exist in evidence pack
Attribution	Reviewer identity + timestamps	reviewer_name / reviewed_at / ai_decision_at present
Exportability	CSV / auto report	Export runs without error
Dispute path	Learner can request human	Dispute fields available when partner delayed

8. Current operating baseline (honest status)

Baseline auto run (no UI screenshots attached to this draft):

Item	Value
Run date (UTC)	17 September 2026, 06:51
Environment	Production data via partner login
Partner account used	nbokete@data-sentinels.com
Evidence folder	docs/aims/evidence/hitl/20260917-065134/
Submissions loaded	3
AI completed	3
Awaiting human (ai_decision empty)	3
Accept / Edit / Reject recorded	0

Interpretation for AIMS:

- Demonstrated: advisory AI grades land; export path works (S01, S07).
- Not yet demonstrated in this baseline: criteria ack on decided rows, dwell on Accept, Edit, Reject, dispute path (S02-S06, S09) - because no partner has completed the new gate decisions on these rows yet.
- Logged gap: one legacy row is approved with a human reviewer (Nono, 26 August 2026) but without ai_decision. This is recorded honestly as pre-gate-field practice, not as proof that AI approved alone. Future approvals should use accept / edit / reject.

Statement for this draft: The gate control is designed and implemented in product. The operational evidence pack for full HITL (Accept/Edit/Reject with dwell) is in progress. This document defines the standard; screenshots and completed S02-S06 runs will be annexed when executed.

9. Evidence annex plan (screenshots later - not in this draft)

When UI evidence is collected, attach under Drive 04 Gate Records / T4L-grading / YYYY-MM-DD/ using names:

Annex ID	Content	Maps to
A1	Terminal / auto report baseline	S01, S07
A2	CSV headers + sample rows	S07, monitoring feed
A3	Criteria blocked to acknowledged	S02
A4	Accept blocked by dwell	S03
A5	Accept after dwell saved	S04
A6	Edit with score delta	S05
A7	Reject saved	S06
A8	Learner explain + dispute (optional)	S09
A9	Points not AI-awarded before/after	S10

Until annexes exist, this document remains valid as the procedure and control definition.

10. Escalation and concerns

Trigger	Action
Partner suspects unfair AI score	Use Edit or Reject; note reason in partner notes
Repeated rubber-stamp pattern (very low dwell)	Raise in weekly delivery sync; review HITL export
AI outage / mass errors	Use manual scoring; log as supplier incident (R-T4L-07)
Privacy concern (wrong tool / leak)	Stop; escalate; follow Tier 3 data rules
Gate control weakened in a release (dwell removed, etc.)	Stop expansion; update Impact + Risk + this gate record same day

Concerns should be raised early - logging a gap is expected; hiding it fails the audit culture AIMS described.

11. Cadence (link to Section 4 Monitoring / Section 7)

Cadence	Gate-related activity
Per decision	Partner follows Section 5; system writes Section 6 fields
Weekly delivery sync	Blockers on backlog / awaiting-human counts
Monthly	Review Accept / Edit / Reject mix, average dwell, score deltas from CSV (feeds Section 4 Monitoring)
On change	Re-test gate after model / prompt / rubric / UI gate changes

12. Open actions before calling the gate "audit-ready"

#	Action	Owner	Status
1	Complete UI Accept path with >=25s dwell on a fresh AI-graded row	Partner + Syntiche	Open
2	Complete Edit path with human != AI score	Partner + Syntiche	Open
3	Complete Reject path	Partner + Syntiche	Open

#	Action	Owner	Status
4	Re-run hitl_gate_test.py --auto and attach updated CSV/report	Syntiche	Open
5	Attach screenshot annex A1-A7 (real dates)	Syntiche	Open
6	Decide partner SLA for awaiting-human backlog (R-T4L-04)	Delivery + Syntiche	Open
7	Record legacy approve-without-ai_decision as known historical gap	Syntiche	Logged (baseline 17 Sep 2026)

13. Statement of gate effectiveness (draft)

Design effectiveness: Adequate for Medium residual risk when operated as designed (aligned to Impact + Risk drafts).

Operating effectiveness (as of 17 September 2026): Partial.

- Design controls are in production code.
- Baseline export proves advisory grades and evidence schema.
- Full operating proof (Accept/Edit/Reject + dwell + criteria on decided rows) is not yet complete and must not be claimed as complete until Section 12 actions are closed.

This partial status is intentional and honest for Stage 1 readiness culture.

14. Sign-off

Role	Name	Signature	Date
Author / System owner	Syntiche Musawu		17 September 2026
Technical / product review			
AIMS coordinator review			
Management acknowledgement			

Author acknowledgement: I confirm this draft defines the T4L grading human gate as implemented, links to Impact/Risk controls, records the 17 September 2026 baseline honestly, and does not claim screenshot / full HITL operating evidence that has not yet been completed.

Next review: after first complete Accept/Edit/Reject evidence pack, or within one month, whichever is sooner.

Appendix A - One-page auditor path

1. Read this gate definition (Sections 2-5).
2. Open linked Risk rows R-T4L-01 / 03 / 04.
3. Inspect baseline folder 20260917-065134 (CSV + auto report).
4. Ask for annex screenshots A3-A7 once marked complete.
5. Confirm monthly monitoring will consume the same CSV fields.

One-sentence control statement: T4L grading may advise; only a partner's dated accept, edit, or reject - after criteria acknowledgement and (for accept) minimum dwell - authorises treating the outcome as human-final, and AI never awards programme points alone.

End of DRAFT - Section 3 Gate / HITL Control Record - T4L grading (no screenshots in this version)